

```

from google.colab import drive

from tensorflow.keras import layers, Sequential, Input, Model
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.preprocessing import image_dataset_from_directory

import numpy as np
import matplotlib.pyplot as plt

from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
from torchvision.transforms import Resize, Normalize, ToTensor, Compose, RandomVerticalFlip, RandomHorizontalFlip, RandomPerspective, RandomInvert, RandomCrop

drive.mount('/content/drive')
dataset_path = '/content/drive/MyDrive/fake_image_detection'

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

```

def get_dataset_split_tensorflow(directory, image_size=(256,256), batch_size=32):
    train_data = image_dataset_from_directory(directory=f'{dataset_path}/train', label_mode='binary', image_size=image_size, batch_size=batch_size)
    val_data = image_dataset_from_directory(directory=f'{dataset_path}/valid', label_mode='binary', image_size=image_size, batch_size=batch_size)
    test_data = image_dataset_from_directory(directory=f'{dataset_path}/test', label_mode='binary', image_size=image_size, batch_size=batch_size)

    return train_data, val_data, test_data

```

```

train_data_tensorflow, val_data_tensorflow, test_data_tensorflow = get_dataset_split_tensorflow(dataset_path)

Found 189 files belonging to 2 classes.
Found 21 files belonging to 2 classes.
Found 90 files belonging to 2 classes.

```

```

data_augmentation = Sequential([layers.RandomFlip("horizontal and vertical"), layers.RandomTranslation(0.3,0.3), layers.RandomRotation(0.4), layers.Rar

```

```

from keras import backend as K

def get_recall(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    recall = true_positives / (possible_positives + K.epsilon())
    return recall

def get_precision(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def get_f1(y_true, y_pred):
    precision = get_precision(y_true, y_pred)
    recall = get_recall(y_true, y_pred)
    return 2*((precision*recall)/(precision+recall+K.epsilon()))

```

```

from keras import initializers

def CNNmodel(input_shape):
    inputs = Input(shape=input_shape+(3,))
    x = data_augmentation(inputs)

    # 1st Convolutional Layer
    x = layers.Conv2D(16, 3, strides = 2, kernel_initializer = 'he_uniform',
    bias_initializer = initializers.Zeros(), padding = 'same')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(3, strides = 1, padding = 'same')(x)

    # 2nd Convolutional Layer
    x = layers.Conv2D(32, 3, strides = 2, kernel_initializer = 'he_uniform',
    bias_initializer = initializers.Zeros(), padding = 'same')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(3, strides = 1, padding = 'same')(x)

    # 3rd Convolutional Layer
    x = layers.Conv2D(64, 3, strides = 2, kernel_initializer = 'he_uniform',
    bias_initializer = initializers.Zeros(), padding = 'same')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(3, strides = 1, padding = 'same')(x)

    # 4th Convolutional Layer
    x = layers.Conv2D(64, 3, strides = 2, kernel_initializer = 'he_uniform',
    bias_initializer = initializers.Zeros(), padding = 'same')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(3, strides = 1, padding='same')(x)

```

```
# Passing it to a Fully Connected layer
x = layers.Flatten()(x)

# 1st Fully Connected Layer
x = layers.Dense(256, kernel_initializer = 'he_uniform',
bias_initializer=initializers.Zeros()(x))
x = layers.BatchNormalization()(x)
x = layers.Activation('relu')(x)

# Output Layer
outputs = layers.Dense(1, activation="sigmoid")(x)
return Model(inputs, outputs)
```

```
model = CNNModel(input_shape=(256,256))
model.summary()
```

sequential_2 (Sequential)	(None, 256, 256, 3)	0
conv2d_60 (Conv2D)	(None, 128, 128, 16)	448
batch_normalization_75 (Batch Normalization)	(None, 128, 128, 16)	64
activation_75 (Activation)	(None, 128, 128, 16)	0
max_pooling2d_60 (MaxPooling2D)	(None, 128, 128, 16)	0
conv2d_61 (Conv2D)	(None, 64, 64, 32)	4640
batch_normalization_76 (Batch Normalization)	(None, 64, 64, 32)	128
activation_76 (Activation)	(None, 64, 64, 32)	0
max_pooling2d_61 (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_62 (Conv2D)	(None, 32, 32, 64)	18496
batch_normalization_77 (Batch Normalization)	(None, 32, 32, 64)	256
activation_77 (Activation)	(None, 32, 32, 64)	0
max_pooling2d_62 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_63 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_78 (Batch Normalization)	(None, 16, 16, 64)	256
activation_78 (Activation)	(None, 16, 16, 64)	0
max_pooling2d_63 (MaxPooling2D)	(None, 16, 16, 64)	0
flatten_17 (Flatten)	(None, 16384)	0
dense_38 (Dense)	(None, 256)	4194560
batch_normalization_79 (Batch Normalization)	(None, 256)	1024
activation_79 (Activation)	(None, 256)	0
dense_39 (Dense)	(None, 1)	257

=====
 Total params: 4,257,057
 Trainable params: 4,256,193
 Non-trainable params: 864

```
from tensorflow.keras import optimizers
```

```
epochs = 1
optimizer = optimizers.Adam()
loss = "binary_crossentropy"
```

```
model.compile(optimizer=optimizer, loss=loss, metrics=['accuracy', get_f1, get_precision, get_recall])
history = model.fit(train_data_tensorflow, epochs=epochs, validation_data=val_data_tensorflow)
```

```
6/6 [=====] - 25s 3s/step - loss: 1.6829 - accuracy: 0.4339 - get_f1: 0.4722 - get_precision: 0.4567 - get_recall: 0.518
```

```
test_loss, test_acc, test_f1, test_precision, test_recall = model.evaluate(test_data_tensorflow)
print('\nTest accuracy:', round(test_acc*100,4))
print('Test recall:', round(test_recall*100,4))
print('Test f1_score:', round(test_f1*100,4))
```

```
3/3 [=====] - 1s 226ms/step - loss: 4.8495 - accuracy: 0.5222 - get_f1: 0.6560 - get_precision: 0.5035 - get_recall: 0.9
```

Test accuracy: 52.2222
Test recall: 98.4127
Test f1_score: 65.6022

2. Tuning the CNN model

```
from tensorflow.keras import applications, regularizers

def ResNet(pre_trained, input_shape):
    inputs = Input(shape=input_shape+(3,))
    x = data_augmentation(inputs)
    x = pre_trained(x)
    x = layers.Flatten()(x)

    x = layers.Dropout(0.5)(x)
    x = layers.Dense(1024, activation="relu")(x)
    x = layers.Dense(512, activation="relu")(x)
    x = layers.Dense(128, activation="relu")(x)
    output = layers.Dense(1, activation='sigmoid')(x)
    return Model(inputs, output)

pre_trained = applications.ResNet50(
    include_top=False,
    weights="imagenet",
    input_shape=(256,256,3),
)

pre_trained.trainable = False

tuned_model = ResNet(pre_trained, input_shape=(256,256))
tuned_model.summary()
```

Model: "model_18"

Layer (type)	Output Shape	Param #
=====		
input_25 (InputLayer)	[(None, 256, 256, 3)]	0
sequential_2 (Sequential)	(None, 256, 256, 3)	0
resnet50 (Functional)	(None, 8, 8, 2048)	23587712
flatten_18 (Flatten)	(None, 131072)	0
dropout_2 (Dropout)	(None, 131072)	0
dense_40 (Dense)	(None, 1024)	134218752
dense_41 (Dense)	(None, 512)	524800
dense_42 (Dense)	(None, 128)	65664
dense_43 (Dense)	(None, 1)	129
=====		
Total params: 158,397,057		
Trainable params: 134,809,345		
Non-trainable params: 23,587,712		

```
from tensorflow.keras import optimizers
from tensorflow.keras import losses

epochs = 10
optimizer = optimizers.Adam(learning_rate=0.01)
loss = losses.BinaryCrossentropy(from_logits=False)

tuned_model.compile(optimizer=optimizer, loss=loss, metrics=['acc', get_f1, get_precision, get_recall])
history = tuned_model.fit(train_data_tensorflow, epochs=epochs, validation_data=val_data_tensorflow)
```

```
Epoch 1/10
6/6 [=====] - 78s 12s/step - loss: 1759.3196 - acc: 0.4233 - get_f1: 0.3604 - get_precision: 0.3427 - get_recall: 0.5577
Epoch 2/10
6/6 [=====] - 72s 12s/step - loss: 5.5496 - acc: 0.4392 - get_f1: 0.4108 - get_precision: 0.3055 - get_recall: 0.6667 -
Epoch 3/10
6/6 [=====] - 77s 13s/step - loss: 4.7985 - acc: 0.5079 - get_f1: 0.3371 - get_precision: 0.2552 - get_recall: 0.5000 -
Epoch 4/10
6/6 [=====] - 69s 12s/step - loss: 4.7167 - acc: 0.6085 - get_f1: 0.3764 - get_precision: 0.3044 - get_recall: 0.5000 -
Epoch 5/10
6/6 [=====] - 69s 12s/step - loss: 3.0805 - acc: 0.4180 - get_f1: 0.2943 - get_precision: 0.2091 - get_recall: 0.5000 -
Epoch 6/10
6/6 [=====] - 73s 12s/step - loss: 1.3664 - acc: 0.4921 - get_f1: 0.3306 - get_precision: 0.2471 - get_recall: 0.5000 -
Epoch 7/10
6/6 [=====] - 75s 13s/step - loss: 0.9957 - acc: 0.4709 - get_f1: 0.2099 - get_precision: 0.1539 - get_recall: 0.3333 -
Epoch 8/10
6/6 [=====] - 69s 12s/step - loss: 0.9109 - acc: 0.4868 - get_f1: 0.3266 - get_precision: 0.2448 - get_recall: 0.5000 -
Epoch 9/10
6/6 [=====] - 73s 12s/step - loss: 0.7870 - acc: 0.5079 - get_f1: 0.4481 - get_precision: 0.3385 - get_recall: 0.6667 -
Epoch 10/10
6/6 [=====] - 73s 12s/step - loss: 0.7720 - acc: 0.4815 - get_f1: 0.2117 - get_precision: 0.1570 - get_recall: 0.3333 -
```

```

test_loss_tuned, test_acc_tuned, test_f1_tuned, test_precision_tuned, test_recall_tuned = tuned_model.evaluate(test_data_tensorflow)
print('\nTest accuracy:', round(test_acc_tuned*100,4))
print('Test recall:', round(test_recall_tuned*100,4))
print('Test f1_score:', round(test_f1_tuned*100,4))

```

```

3/3 [=====] - 26s 9s/step - loss: 0.7693 - acc: 0.5000 - get_f1: 0.6651 - get_precision: 0.5000 - get_recall: 1.0000

```

```

Test accuracy: 50.0
Test recall: 100.0
Test f1_score: 66.5121

```

```

from keras import optimizers, regularizers
from keras.callbacks import EarlyStopping

```

```

early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    verbose=1,
    mode='min',
    restore_best_weights=True
)

```

```

learning_rate = 0.0001
batch_size = 32
num_epoch = 5
optimizers = [optimizers.SGD(), optimizers.Adam()]
regularizers = [regularizers.l1(0.01), regularizers.l2(0.01), regularizers.l1_l2(l1=0.01, l2=0.01)]

```

```

for opt in optimizers:
    for reg in regularizers:
        model = CNNModel(input_shape = (256, 256))
        model.compile(
            optimizer = opt,
            loss = 'binary_crossentropy',
            metrics = ['accuracy']
        )

        history = model.fit(
            train_data_tensorflow,
            epochs = num_epoch,
            batch_size = batch_size,
            validation_data = val_data_tensorflow,
            callbacks = [early_stopping],
            verbose = 1
        )

        # Evaluate the model
        loss, acc = model.evaluate(test_data_tensorflow)
        print(f'\nOptimizer: {opt}, \nRegularizer: {reg}, \n loss: {loss}, \n accuracy: {acc}\n')

```

Epoch 1/5
6/6 [=====] - 15s 2s/step - loss: 1.6891 - accuracy: 0.4921 - val_loss: 7.5799 - val_accuracy: 0.5238
Epoch 2/5
6/6 [=====] - 13s 2s/step - loss: 0.9084 - accuracy: 0.4709 - val_loss: 2.0537 - val_accuracy: 0.5238
Epoch 3/5
6/6 [=====] - 13s 2s/step - loss: 0.6636 - accuracy: 0.5661 - val_loss: 7.6989 - val_accuracy: 0.4762
Epoch 4/5
6/6 [=====] - 18s 3s/step - loss: 0.7550 - accuracy: 0.5397 - val_loss: 8.0562 - val_accuracy: 0.4762
Epoch 5/5
6/6 [=====] - 13s 2s/step - loss: 0.7145 - accuracy: 0.5503 - val_loss: 5.2279 - val_accuracy: 0.4762
3/3 [=====] - 1s 226ms/step - loss: 6.0498 - accuracy: 0.4889

Optimizer: <keras.optimizers.legacy.adam.Adam object at 0x7f8c4409d1c0>,
Regularizer: <keras.regularizers.L1L2 object at 0x7f8c4409d250>,
loss: 6.0497565269470215,
accuracy: 0.488888895511627

Optimizer: optimizers.Adam()
Regularizer: regularizers.L1_L2(l1=0.01, l2=0.01)
Loss: 1.2607940435409546
Accuracy: 0.5444444417953491